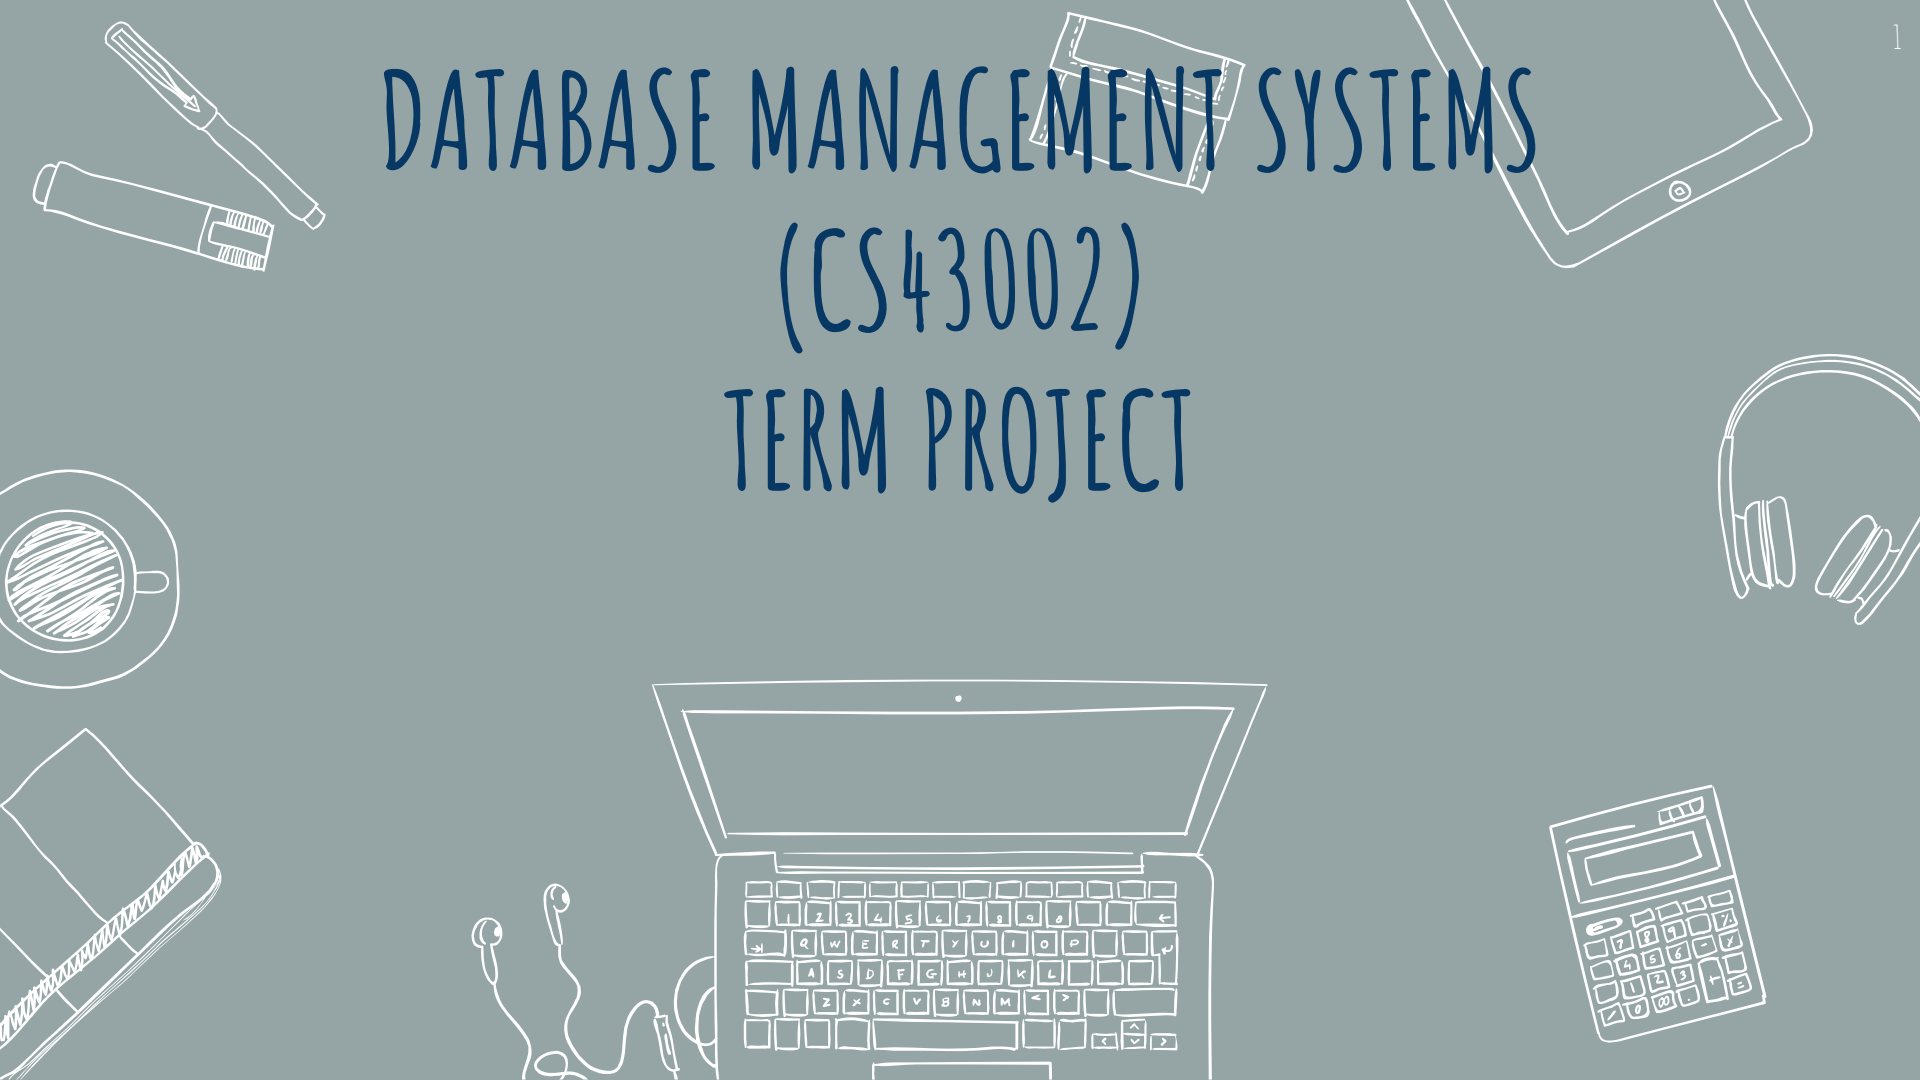


# DATABASE MANAGEMENT SYSTEMS

## (CS43002)

### TERM PROJECT





# QUERY PERFORMANCE ANALYSIS AND OPTIMIZATION THROUGH REAL-TIME METRICS COLLECTION IN POSTGRES

Group members (AARVY):

1. Aditya Choudhary (20CS10005)
2. Astitva (20CS30007)
3. Rishi Raj (20CS30040)
4. Vikas Vijaykumar Bastewad (20CS10073)
5. Yashraj Singh (20CS10079)



# WHY DO WE NEED THIS ?

- Finding bottlenecks in a query processing system.
- Making system reliable.
- Building an alerting mechanism to prevent database damage.
- Analyze query performance and optimize resources.



# PREVIOUS WORK IN THIS DOMAIN

- `pg_stat_statements` module in Postgres, which collects statistics on the execution of SQL statements.
- `Sys.dm_exec_query_stats` in Microsoft SQL Server, which collects query performance metrics such as execution count, total execution time, and query plan analysis
- None of them provide real-time metrics collection during query execution.

# OUR SOLUTION

- Build wrapper that can collect and report query processing metrics in real time.
- The collected metrics will be reported in a user-friendly format

# INBUILT-POSTGRES STATISTICS

**PostgreSQL offers a built-in statistics collector that automatically aggregates key metrics.**

- `pg_stat_database`
- `pg_stat_user_tables`
- `pg_stat_user_indexes`
- `pg_stat_bgwriter`
- `pg_statio_user_tables`





# QUERYPILOTX STATISTICS

- **get\_stats\_disk\_usage\_for\_database:**
  - Gives the size of each database in bytes
- **get\_stats\_index\_hit\_rates:**
  - Gives indexes vs sequential scan through the table, helping us understand effectiveness of our indexing strategies.
- **get\_stats\_tx\_rate\_for\_database:**
  - Gives the rate of transactions executed per second as well as the rate of roll backs executed per second.
- **get\_stats\_seconds\_since\_last\_vacuum\_per\_table:**
  - Gives the amount of time passed in seconds since the last vacuum was performed on each table in the database.



# SOME REAL-TIME METRICS

- **Query Duration:**
  - Reveals the time duration for which the query has been running.
- **CPU usage:**
  - Gives the percentage of CPU time consumed by each active query.
- **Memory Usage:**
  - Reveals the amount of memory consumed by each active query.
- **Disk I/O:**
  - Shows the number of bytes being read or written by each active query.





# INTERFACES DEVELOPED BY US

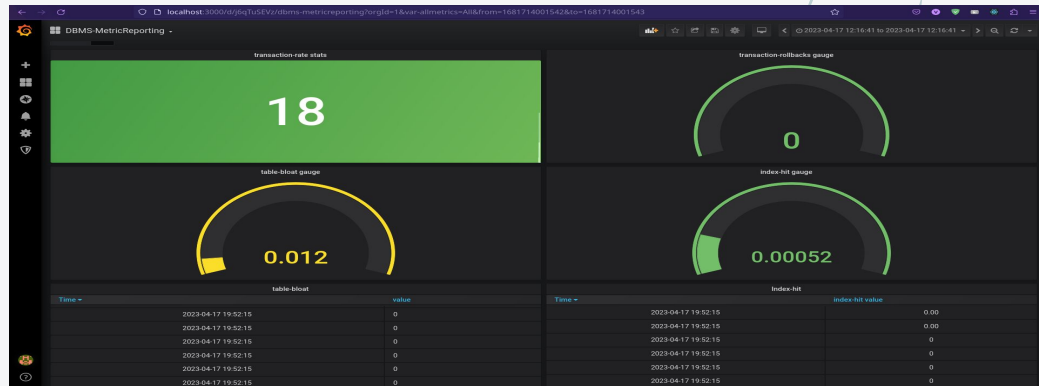
We have made three tools to provide different kinds of analytics

1. **Python CLI's app integrated with Grafana**
2. **Python GUI App**
3. **Go CLI app to provide in-depth table and overall database analysis**

# PYTHON CLI APP INTEGRATED WITH GRAFANA

Python is used to connect to the database and fetch different metrics. These fetched raw data are then written in a log file, which can then be used by users to integrate with different monitoring tools. currently, we have integrated it with Grafana.

```
term-project/postgresql_metrics/postgresql_metrics via @ v3.10.10 took 5s
→ python3 metrics logic.py all
postgresql-metrics.yml
[2023-04-16 14:34:42.335172] INFO: postgresql-metrics: open database connection to 127.0.0.1:5432, user 'postgres', database 'test'
# sleep 5 s to get diffs on derivative metrics
{'type': 'metric', 'key': 'postgresql', 'value': 7, 'attributes': {'what': 'client-connections', 'unit': 'connection'}}
{'type': 'metric', 'key': 'postgresql', 'value': 1, 'attributes': {'what': 'locks_granted', 'type': 'locks', 'locktype': 'relation', 'unit': 'lock'}}
{'type': 'metric', 'key': 'postgresql', 'value': 0, 'attributes': {'what': 'locks_waiting', 'type': 'locks', 'locktype': 'relation', 'unit': 'lock'}}
{'type': 'metric', 'key': 'postgresql', 'value': 1, 'attributes': {'what': 'locks_granted', 'type': 'locks', 'locktype': 'virtualxid', 'unit': 'lock'}}
{'type': 'metric', 'key': 'postgresql', 'value': 0, 'attributes': {'what': 'locks_waiting', 'type': 'locks', 'locktype': 'virtualxid', 'unit': 'lock'}}
{'type': 'metric', 'key': 'postgresql', 'value': 2, 'attributes': {'what': 'locks_granted', 'type': 'locks', 'locktype': 'total', 'unit': 'lock'}}
{'type': 'metric', 'key': 'postgresql', 'value': 0, 'attributes': {'what': 'locks_waiting', 'type': 'locks', 'locktype': 'total', 'unit': 'lock'}}
{'type': 'metric', 'key': 'postgresql', 'value': 0, 'attributes': {'what': 'wal-file-amount', 'type': 'locks', 'locktype': 'total', 'unit': 'lock'}}
{'type': 'metric', 'key': 'postgresql', 'value': 100.0, 'attributes': {'what': 'xid-remaining', 'unit': '%'}}
{'type': 'metric', 'key': 'postgresql', 'value': 7738159, 'attributes': {'what': 'database-size', 'database': 'test', 'unit': 'B'}}
{'type': 'metric', 'key': 'postgresql', 'value': 2.8, 'attributes': {'what': 'transaction-rate', 'type': 'transactions', 'database': 'test', 'unit': 'transaction'}}
{'type': 'metric', 'key': 'postgresql', 'value': 0.0, 'attributes': {'what': 'transaction-rollback', 'type': 'transactions', 'database': 'test', 'unit': 'transaction'}}
{'type': 'metric', 'key': 'postgresql', 'value': 0, 'attributes': {'what': 'sec-since-oldest-xact-start', 'database': 'test', 'unit': 's'}}
```



# PYTHON GUI APPLICATION

We have made an app that provides real-time query analysis and system resource usage in GUI.

PostgreSQL 15.2 - yashrajs-zenbook-ubuntu - postgres@127.0.0.1:5432/postgres - Ref.: 2s

4.79M - 0B/s - TPS: 1 - Active connections: 4 - Duration mode: query

## RUNNING QUERIES

DATABASE	APP	USER	CLIENT	TIME+	Waiting	state	Query
			local	00:06.09		active	select * from
							, "Album", "Customer", "Invoice";
			local	00:05.14		active	select * from
							Line, "MediaType", "Track", "Playlist";
			local	00:04.19		active	select * from "Track"
							"Employee", "Genre";

PostgreSQL 15.2 - rsh-raj -

\* Global: 1 day, 22 hours and 20 minutes uptime, 38.81M dbs

Sessions: 2/100 total, 2 active, 0 idle, 0 idle in txn, 0 idle in txn abrt, 0 waiting

Activity: 3 tps, 0 insert/s, 0 update/s, 0 delete/s, 0 tuples returned/s, 0 temp files, 0B temp size

\* Worker processes: 0/8 total, 0/4 logical workers, 0/8

Other processes & info: 0/3 autovacuum workers, 0/10 wal

\* Mem.: 7.60G total, 214.51M (2.76%) free, 5.95G (78.32%) used, 1.44G (18.92%) buff+cached

Swap: 22.79G total, 21.09G (92.54%) free, 1.70G (7.46%) used

IO: 0/s max iops, 0B/s - 0/s read, 0B/s - 0/s write

Load average: 0.07 0.49 0.66

## RUNNING QUERIES

PID	DATABASE	APP	USER				
219186				local	97.2	0.2	0B 0B 00:03.66

# GO CLI APPLICATION

This app provides in-depth table and overall database analysis in the form of two reports, whole DB analysis, and individual table analysis.

```
PostgreSQL Cluster:
Name: 15/main
Server Version: 15.2 (Ubuntu 15.2-1.pgdg22.04+1)
Server Started: 10 Apr 2023 5:55:01 PM (6 days ago)
System Identifier: 7194855233783590104
Timeline: 1
Last Checkpoint: 10 Apr 2023 5:55:01 PM (6 days ago)
REDO LSN: 0/2433AD0
Checkpoint LSN: 0/2433AD0 (0 B since REDO)
Transaction IDs: oldest = 716, next = 1165, range = 449
Notification Queue: 0.0% used
Active Backends: 1 (max 100)
Recovery Mode? no
```

```
System Information:
Hostname: yashrajs-zenbook-ubuntu
CPU Cores: 16 x Intel(R) Core(TM) i9-10980HK CPU @ 2.40GHz
Load Average: 0.37
Memory: used=4.6 GiB, free=11 GiB, buff=1.4 GiB, cache=13 GiB
Swap: used=0 B, free=3.8 GiB
```

Setting	Value
shared_buffers	16384 (128 MiB)
work_mem	4096 (4.0 MiB)
maintenance_work_mem	65536 (64 MiB)
temp_buffers	1024 (8.0 MiB)
autovacuum_work_mem	-1
temp_file_limit	-1
max_worker_processes	8
autovacuum_max_workers	3
max_parallel_workers_per_gather	2
effective_io_concurrency	1

```
Database #4:
Name: hms
Owner: postgres
Tablespace: pg_default
Connections: 1 (no max limit)
Frozen Xid Age: 449
Transactions: 1621 (100.0%) commits, 0 (0.0%) rollbacks
Cache Hits: 99.8%
Rows Changed: ins 0.0%, upd 0.0%, del 0.0%
Total Temp: 0 B in 0 files
Problems: 0 deadlocks, 0 conflicts
Totals Since:
Size: 7.8 MiB
Installed Extensions:
```

Name	Version	Comment
plpgsql	1.0	PL/pgSQL procedural language

```
Table #1 in "hms":
Name: hms.public.users
Columns: 8
Manual Vacuums: never
Manual Analyze: never
Auto Vacuums: never
Auto Analyze: never
Post-Analyze: 0.0% est. rows modified
Row Estimate: 0.0% live of total 0
Rows Changed: ins 0.0%, upd 0.0%, del 0.0%
HOT Updates: 0.0% of all updates
Seq Scans: 0, 0.0 rows/scan
Idx Scans: 0, 0.0 rows/scan
Cache Hits: 0.0% (idx=0.0%)
Size: 16 KiB
```

Index	Type	Size	Bloat	Cache Hits	Scans	Rows Read/Scan	Rows Fetched/Scan
users_pkey	btree	16 KiB			0	0.0	0.0

```
Table #2 in "hms":
Name: hms.public.doctors
Columns: 4
```

THANK YOU

